

STA323 Assignment1 report

SID: 12110821

Name: ZHANG Chi

Solution for Q1

(1)

There are 8 columns in the dataset, and the missing value is not allowed when the data is read in. There are total 541909 rows.

The schema of the dataset is defined as follows:

```
1  schema = StructType([
2      StructField("InvoiceNo", IntegerType(), False),      # 第三个: 是否允许有空值
3      StructField("StockCode", IntegerType(), False),
4      StructField("Description", StringType(), False),
5      StructField("Quantity", IntegerType(), False),
6      StructField("InvoiceDate", StringType(), False),
7      StructField("UnitPrice", FloatType(), False),
8      StructField("CustomerID", IntegerType(), False),
9      StructField("Country", StringType(), False)
10 ])
11 df = spark.read.csv("./data/Q1_data/retail-dataset.csv", header=True,
12                      schema=schema)
```

The first 5 rows of the dataset are shown below:

```
+-----+-----+-----+-----+-----+-----+-----+
|InvoiceNo|StockCode|Description|Quantity|InvoiceDate|UnitPrice|CustomerID|Country|
+-----+-----+-----+-----+-----+-----+-----+
| 536365| NULL|WHITE HANGING HEA...| 6|12/1/2010 8:26| 2.55| 17850|United Kingdom|
| 536365| 71053| WHITE METAL LANTERN| 6|12/1/2010 8:26| 3.39| 17850|United Kingdom|
| 536365| NULL|CREAM CUPID HEART...| 8|12/1/2010 8:26| 2.75| 17850|United Kingdom|
| 536365| NULL|KNITTED UNION FLA...| 6|12/1/2010 8:26| 3.39| 17850|United Kingdom|
| 536365| NULL|RED WOOLLY HOTTIE...| 6|12/1/2010 8:26| 3.39| 17850|United Kingdom|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

After using `filter` to drop the records where their `quantity` or `UnitPrice` is not positive, 363087 rows are left.

```
1  df_clean = df.dropna().filter((col("Quantity") > 0) & (col("UnitPrice") > 0))
```

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850	United Kingdom
536365	22752	SET 7 BABUSHKA NE...	2	12/1/2010 8:26	7.65	17850	United Kingdom
536365	21730	GLASS STAR FROSTE...	6	12/1/2010 8:26	4.25	17850	United Kingdom
536366	22633	HAND WARMER UNION...	6	12/1/2010 8:28	1.85	17850	United Kingdom
536366	22632	HAND WARMER RED P...	6	12/1/2010 8:28	1.85	17850	United Kingdom

only showing top 5 rows

(2)

The revenue is calculated by multiplying the `Quantity` and `UnitPrice` for each record, and then summing up the results.

```
1 df_clean.select(expr("sum(UnitPrice * Quantity) as total_cost")).show()
```

total_cost
8015349.50373831

(3)

Before get top 5 customers that spend most, we need to calculate the total cost for each customer. Then we can sort the total cost in descending order.

Here I use `groupBy` and `agg` to calculate the total cost for each customer, and then use `orderBy` to sort the result.

```
1 df_clean.groupBy("CustomerID").agg(expr("sum(UnitPrice * Quantity) as
total_cost")).select(col("CustomerID"),col("total_cost")).orderBy("total_cost",a
scending = False).show(5)
```

CustomerID	total_cost
14646	265106.91930553317
18102	253922.7600557804
17450	180847.0303592682
16446	168472.49374997616
14911	125544.34975004196


only showing top 5 rows

(4)

I use `to_timestamp` to convert the `InvoiceDate` to date type. After selecting (by `where`) the data during designated days, I use `groupBy` and `agg` to calculate the total cost for each day.

```
1 # set the timeParserPolicy to LEGACY
2 spark.conf.set("spark.sql.legacy.timeParserPolicy", "LEGACY")
3 df_4 = df_clean.withColumn("InvoiceDate", to_timestamp(col("InvoiceDate"),
4 "MM/dd/yyyy HH:mm"))
5 df_4.where("InvoiceDate between '2010-12-01' and '2010-12-05'")
6   .groupBy("Country").agg(expr("sum(UnitPrice * Quantity) as total_cost")
7   ).select(col("Country"), col("total_cost")).collect()
```

The bar plot drawn by `seaborn` is shown below. To make it clear, I also draw the plot in `log` scale.



Solution for Q2

(1)

I use `os.listdir` to get the list of all CSV files in the directory. Then I read the first CSV file to initialize the DataFrame, and use `union` to merge the rest of the CSV files. Observing that the first row of each CSV file features the double quote as it beginning, to remove it, the `comment` option is used while loading `.csv` files.

```

1 import os
2 csv_list = [p for p in os.listdir(path = "data/Q2_data") if p.endswith(".csv")]
3
4 # 读取第一个 CSV 文件来初始化 DataFrame
5 df =
    spark.read.format("csv").option("header", "True").option("comment", "\\").load(f"
    data/Q2_data/{csv_list[0]}")
6
7 # 从第二个文件开始, 读取每一个 CSV 文件并进行合并
8 for i in csv_list[1:]:
9     temp_df =
        spark.read.format("csv").option("header", "True").option("comment", "\\").load(f"
        data/Q2_data/{i}")
10    df = df.union(temp_df)

```

Before using `spark.read` to load files and remove first lines by `comment`, I tried to use `spark.read.text` and extract the lines by `df.tail()` (except first line) to `createDataFrame`. However, I can't find a easy way to split the each row to make the the number of result correpd to the number of columns as some values in one row is wrapped by "{}" with many commas(,) inside. The graphy is shown below. I can't split the row by a single comma.

```

,13.0,447.456,,66.059,38S294M229S2N,,332S1N37M192S,1.312e-127,,7.106e-
15,98.639,,97.297,39.0,332.0,1.0,294.0,,,,,333.0,369.0,2.0,38.0,39.0,113.0,114.0,137.0,138.0,188.0,189.0,1
97.0,198.0,305.0,339.0,368.0,306.0,338.0,,0.0,,,GGCAGCCCAAGGCTGCCCTCGGTCACTCTGTTCCCGCCCTCCTCTGGGGAGCTTAA
AGCCAACAAGGCCACACTGGTGGGTCTCATAAGTGAAATTCTACCCGGGAGCCGTGACAGTGGCCTGGAAGGCAGATAGCAGCCCCGTCAAGGCCGGGAGTGGAGACC
ACCACACCCACCCCAACAAAAAACAAC,1,"{'fw11': {'1 ': 'Q', '2 ': 'S', '3 ': 'V', '4 ': 'L', '5 ': 'T', '6 ':
'Q', '7 ': 'P', '8 ': 'P', '9 ': 'S', '11 ': 'A', '12 ': 'S', '13 ': 'G', '14 ': 'T', '15 ': 'P', '16 ':
'G', '17 ': 'Q', '18 ': 'R', '19 ': 'V', '20 ': 'T', '21 ': 'I', '22 ': 'S', '23 ': 'C', '24 ': 'S', '25
': 'G', '26 ': 'S'}, 'cdr11': {'27 ': 'S', '28 ': 'S', '29 ': 'N', '30 ': 'I', '35 ': 'G', '36 ': 'S', '37
': 'D', '38 ': 'T'}, 'fw12': {'39 ': 'V', '40 ': 'N', '41 ': 'W', '42 ': 'F', '43 ': 'Q', '44 ': 'Q', '45
': 'L', '46 ': 'P', '47 ': 'G', '48 ': 'S', '49 ': 'Col 96', '50 ': 'P', '51 ': 'K', '52 ': 'L', '53 ': 'L',
'54 ': 'I', '55 ': 'Y'}, 'cdr12': {'56 ': 'S', '57 ': 'N', '65 ': 'N'}, 'fw13': {'66 ': 'Q', '67 ': 'R',
'68 ': 'P', '69 ': 'S', '70 ': 'G', '71 ': 'V', '72 ': 'P', '74 ': 'D', '75 ': 'R', '76 ': 'F', '77 ':
'S', '78 ': 'G', '79 ': 'S', '80 ': 'K', '83 ': 'S', '84 ': 'G', '85 ': 'T', '86 ': 'S', '87 ': 'A', '88
': 'S', '89 ': 'L', '90 ': 'A', '91 ': 'I', '92 ': 'S', '93 ': 'G', '94 ': 'L', '95 ': 'Q', '96 ': 'S',
'97 ': 'E', '98 ': 'D', '99 ': 'E', '100 ': 'A', '101 ': 'D', '102 ': 'Y', '103 ': 'Y', '104 ': 'R'},
'cdr13': {'105 ': 'A', '106 ': 'A', '107 ': 'W', '108 ': 'D', '109 ': 'D', '110 ': 'S', '113 ': 'L', '114
': 'N', '115 ': 'G', '116 ': 'W', '117 ': 'V'}, 'fw14': {'118 ': 'V', '119 ': 'G', '120 ': 'G', '121 ':
'G', '122 ': 'T', '123 ': 'K', '124 ': 'L', '125 ': 'T', '126 ': 'V', '127 ': 'L'}}",|||Missing Conserved
Cysteine: 104||

```

So that the code below is not used in the final version.

```

df = spark.read.text("data.csv")
col_name = (df.tail(2)[0]["value"]).split(",")
data = [df.tail(1)[0]["value"].split(",")]
df1 = spark.createDataFrame(data,col_name)

```

Then filter the data to get the result.

```

1 df.select(col("sequence_alignment_aa"),col("cdr1_aa"),col("cdr2_aa"),col("cdr3_a
a")).where((length(col("cdr3_aa"))>=10) & (length(col("cdr3_aa"))<=100 ))
2 df_remove.coalesce(1).write.format("csv").mode("overwrite").save("output/output_
data/q2_pyspark.csv")

```

```

+-----+-----+-----+
|sequence_alignment_aa| cdr1_aa|cdr2_aa|   cdr3_aa|
+-----+-----+-----+
| QSVLTQPPSVSGTPGQR...| DSNIGNNF|   KTS| AAWDDPLNAVL|
| GVPDRFSGSTSGTSASL...|   NULL|   NULL| QSFDSLGGFYV|
| QSALTQPPSASGTPGQR...| RSNIGINT|   SND| DAWDDNLNGPV|
| QSALTQPASMSGSPGQS...| SSDVGASNH|   EVS| YSYAVGVTFV|
| TSLTCGSSTGAVHSGY...| TGAVHSGYY|   STD| LLYFGGIQPLWV|
+-----+-----+-----+
only showing top 5 rows

```

Save it to a single `.csv` file (Actually, it generated 4 files in directory `q2_pyspark.csv`).

```

1 df_remove.coalesce(1).write.format("csv").mode("overwrite").save("output/output_
  data/q2_pyspark.csv")

```

```

q2_pyspark.csv
_SUCCESS
._SUCCESS.crc
.part-00000-9c1a308a-2505-4f2b-8d74-146ed8211d09-c000.csv.crc
part-00000-9c1a308a-2505-4f2b-8d74-146ed8211d09-c000.csv

```

(2)

Run the script by command `bash q2_2.sh` and the output file `q2(2)_sh.csv` is generated. There are 17157 rows (including column names) in the output file `q2(2)_sh.csv`.

```

1 sequence_alignment_aa,cdr1_aa,cdr2_aa,cdr3_aa
2 QSVLTQPPSASGTPGQRTISCSGSSNIGSDTVNWFQQLPGSAPKLLIYSSNQRPSPGVDRFSGSKSGTSASLAISGLQSEDEADYYRAAW
3 DDLSLNGWVVGGGTKLTVL,SSNIGSDT,SSN,AAWDDSLNGWV
4 QSMLTQPPSASGTPGQRTISCSGGNSNIGSNTVSWYQQFPGAAPKLLIYSSNQRPSPGVPARFSGSRSGTSASLAISGLQSEDEAVYYCASW
5 DDGLDGFVIFGAGTKLTVL,NSNIGSNT,SSN,ASWDDGLDGFVI
6 QSMLTQPPSASGTPGQRTISCSGGNSNIGSNTVSWYQQFPGAAPKLLIYSSNQRPSPGVPARFSGSRSGTSASLAISGLQSEDEAVYYCASW
7 DDGLDGFVIFGAGTKLTVL,NSNIGSNT,SSN,ASWDDGLDGFVI
8 QSMLTQPPSASGTPGQRTISCSGGNSNIGSNTVSWYQQFPGAAPKLLIYSSNQRPSPGVPARFSGSRSGTSASLAISGLQSEDEAVYYCASW
9 DDGLDGFVIFGAGTKLTVL,NSNIGSNT,SSN,ASWDDGLDGFVI
10 GVPDRFSGSKSGTSASLAITGLQAEDESAYYCQSYDNLSSVWVFGGGTMTLTVL,,QSYDNLSSVWV
11 QSVLTQPPSASGTPGQRTISCSGSSNIGSNYVYWYQQLPGTAPKLLIYRNNQRPSPGVDRFSGSKSGTSASLAISGLRSEDEADYYCAAW
12 DDLSLNGWVVGGGTKLTVL,SSNIGSNY,RNN,AAWDDSLSGWV

```

(3)

Run the script by command `bash q2_3.sh` and the output file `q2(3)_sh.csv` is generated. There are 17140 rows (including column names) in the output file `q2(3)_sh.csv`.

```

1 sequence_alignment_aa,cdr1_aa,cdr2_aa,cdr3_aa
2 EIVMTQSPATLSVSPGERATLSCRASQSVSSNLAWYQHKPGQAPRLLIYGTSRATGIPARFSGSGSGTEFTLTISSLQSEDF
3 VYYCHQYNSWPPGTFGQGTKEI,QSVSSN,GTS,HQYNSWPPGT
4 DIQMTQSPSSLSASVGRVTITCRASQSISSYLNWYQKPGKAPKLLIYAASSLQSGVPSRFSGSGSGTDFTLTISSLQPEDFA
5 TYQCQQSYSTHPYTFGQGTKEI,QSISSY,AAS,QQSYSTHPYT
6 EIVMTQSPATLSVSPGERATLSCRASQSVSSNLAWYQKPGQAPRLLIYGTSRATGIPARFSGSGSGTEFTLTISSLQSEDF
7 VYYCQQYNNWPPWTFGQGTKEI,QSVSSN,GAS,QQYNNWPPWT
8 EIVLAQSPATLSVSPGERATLSCRASQSVSSYLNWYQKPGQAPRLLIFDASNRATGIPARFSGSGSGTDFTLTISLLEPEDFA
9 VYYCQQRNNWPPYTFGQGTKEI,QSVSSY,DAS,QQRNNWPPYT
10 DIVLTQSPGTLSPGERATLSCRATHSINRRFMAWYRQKGQAPRVIIYGTIRATGIPDRFSGSGSGTDFTLTISRLEAEDS
11 AVYYCQQYDTSQGYFPGQGTKEI,HSINRRF,GTS,QQYDTSQGYF

```