

STA323 Assignment 2 report

SID: 12110821

Name: ZHANG Chi

Solution for Q1

(1)

As the header of each sequence is started with `>`, we can drop it by `filter` after reading the file by `spark.read.text`.

```
1 df1_1 = spark.read.text("data/Q1_data/protein.fasta")
2 df1_1 = df1_1.filter(~col("value").contains(">"))
```

```
+-----+
|value                                     |
+-----+
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|M
|KKEVVAVAKKEEVLKKEVVVPSKKDEEILPLKKEVPRPPKKEEDVMPQKKEVPRPPKKEE|
|DIVPQMRDVSLLPPKEEEEKIVPKKKEVPRPPKKVEEILPPKKEVHRPPKKEEDIVPQIREV|
|SLPPKKDEEIVCEKKEVAPAKEEPSKKPKVPSLPATQREDVIEEIIHKKPTAALSKFEDV|
|KEHEEKETFVVLKKEIIDAPTKKEMVTAKHVIVPQKEEIIIPSPTQEEVVSFKRKQTVRTS|
+-----+
only showing top 5 rows
```

Then we can split each line by `split` and `explode` them to make every char (amino acid) in one line in the column `chars`. Finally, we can use `groupBy` and `count` to get the frequency of each word.

```
1 df1_1_withchars = df1_1.withColumn("chars", explode(split(col("value"), "")))
2 df1_1_withchars.groupBy("chars").count().show(5)
```

```
+-----+-----+
|value                                     |chars|
+-----+-----+
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|M|
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|E|
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|E|
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|I|
|MEEITQIKKRLSQTVRLEGKEDLLSKKDSITNLKTEEHVSVKKMWISEPKPEKKEDIQLK|T|
+-----+-----+
only showing top 5 rows
```

```

+-----+-----+
|chars|  count|
+-----+-----+
|    K|1684031|
|    F| 985877|
|    Q|1422769|
|    E|2674664|
|    T|2795042|
+-----+-----+
only showing top 5 rows

```

(2)

Using RDD api, we can read the file and drop the header by `filter` and `map` the lines to characters. Then we can use `flatMap` to make every char in one line and use `countByValue` to get the frequency of each word.

```

1 rdd1_2 = spark.sparkContext.textFile("data/Q1_data/protein.fasta")
2 rdd1_2 = rdd1_2.filter(lambda x: ">" not in x)
3 rdd1_2_withchars = rdd1_2.flatMap(lambda x: list(x))
4 total_count = rdd1_2_withchars.count()
5
6 frequencies = {k: v / total_count for k, v in counts.items()}
7 frequencies

```

```

{'M': 0.013064028899518616,
'E': 0.07474616297912196,
'I': 0.04826036842051577,
'T': 0.07811024669472166,
'Q': 0.0397607039821235,
'K': 0.047061022111050006
}

```

Besides, we can also use `reduceByKey` to get the count of each word, which could be more friendly to a low memory driver.

```

1 rdd1_2_withchars.map(lambda x: (x, 1)).reduceByKey(lambda x, y: x + y).take(5)
✓ 6.0s

[('M', 467474), ('E', 2674664), ('I', 1726915), ('Q', 1422769), ('K', 1684031)]

```

(3)

To find the a specific sequence motif "STAT" (omitting the line break), we can use `re.findall` in `re` module to find all the matches in the every element of the rdd object `rdd1_2`. After that, we can use `filter` to drop the elements without any match and use `flatMap` to make every match in one line. Finally, we can use `count` to get the number of the matches.

```
1 import re
2 rdd1_2.map(lambda x: list(re.findall("STAT",x))).filter(lambda x: len(x)!=0).flatMap(lambda x: x).count()
✓ 0.4s
2052
```

Solution for Q2

(1)

Firstly, read two `.csv` files by `spark.read.csv`, then rename some columns by `withColumnRenamed` as well as convert the data type of some columns if necessary according to the following questions.

```
1 course = spark.read.csv("data/Q2_data/courses.csv",header=True)
2 course = course.withColumnRenamed("title","course_title")
3 course = course.withColumn("created",to_timestamp("created", "yyyy-MM-dd'T'HH:mm:ss'Z'"))
```

id	course_title	url	rating	num_reviews	num_published_lectures	created	last_update_date	duration	instructors_id	image
567828	The Complete Pyth...	/course/complete-...	4.5927815	452973	155	2015-07-29 00:12:23	2021-03-14	22 total hours	9685726	https://img-c.ude...
1565838	The Complete 2023...	/course/the-compl...	4.667258	263152	490	2018-02-22 12:02:33	2023-01-20	65.5 total hours	31334738	https://img-c.ude...
625204	The Web Developer...	/course/the-web-d...	4.6961474	254711	616	2015-09-28 21:32:19	2023-02-12	64 total hours	4466306	https://img-c.ude...
756150	Angular - The Com...	/course/the-compl...	4.5926924	180257	472	2016-02-08 17:02:55	2023-02-06	34.5 total hours	13952972	https://img-c.ude...
2776760	100 Days of Code...	/course/100-days-...	4.6952515	177568	676	2020-01-24 10:47:21	2022-11-30	64 total hours	31334738	https://img-c.ude...

only showing top 5 rows

```
1 instructors = spark.read.csv("data/Q2_data/instructors.csv",header=True)
2 instructors = instructors.withColumnRenamed("title","instructor_title")
```

class	id	instructor_title	name	display_name	job_title	image_50x50	image_100x100	initials	url
user	9685726	Jose Portilla	Jose	Jose Portilla	Head of Data Scie...	https://img-c.ude...	https://img-c.ude...	JP	/user/joseportilla/
user	31334738	Dr. Angela Yu	Dr. Angela	Dr. Angela Yu	Developer and Lea...	https://img-c.ude...	https://img-c.ude...	DY	/user/4b4368a3-b5...
user	4466306	Colt Steele	Colt	Colt Steele	Developer and Boo...	https://img-b.ude...	https://img-b.ude...	CS	/user/coltsteele/
user	13952972	Maximilian Schwar...	Maximilian	Maximilian Schwar...	AWS certified, Pr...	https://img-b.ude...	https://img-b.ude...	MS	/user/maximilian-...
user	599932	Tim Buchalka	Tim	Tim Buchalka	Java Python Andro...	https://img-c.ude...	https://img-c.ude...	TB	/user/timbuchalka/

Then `join` them together (inner join) by columns `instructor_id` and `course_id`.

```
1 df2_1 = course.join(instructors, course["instructors_id"] == instructors["id"], "inner")
2 print(df2_1.count())
3 df2_1.show(3)
```

✓ 0.5s

id	course_title	url	rating	num_reviews	num_published_lectures	created	last_update_date	duration	instructors_id	image	class
567828	The Complete Pyth...	/course/complete-...	4.5927815	452973	155	2015-07-29 00:12:23	2021-03-14	22 total hours	9685726	https://img-c.ude...	user
1565838	The Complete 2023...	/course/the-compl...	4.667258	263152	490	2018-02-22 12:02:33	2023-01-20	65.5 total hours	31334738	https://img-c.ude...	user
625204	The Web Developer...	/course/the-web-d...	4.6961474	254711	616	2015-09-28 21:32:19	2023-02-12	64 total hours	4466306	https://img-c.ude...	user

only showing top 3 rows

(2)

Before selecting by `spark SQL`, a temporary view should be created by `createOrReplaceTempView`.

Here is the detail about the SQL query referring to the requirements:

- **Highest course rating:** order the data by `rating` in descending order and select the first row.
- **Among all courses that are related to 'spark':** use `like` in `where` to filter the rows that contain 'spark' in the column `course_title`.

- **Created after 2018-01-01 00:00:00:** select the rows that the column `created` is later than `2018-01-01 00:00:00` in `where` .

```
1 df2_1.createOrReplaceTempView("df2_1")
2 spark.sql("select display_name, job_title from df2_1 where course_title like
'%spark%' and created > '2018-01-01 00:00:00' order by rating desc LIMIT
1").show(truncate=False)
```

```
+-----+-----+
|display_name|job_title|
+-----+-----+
|Deby Coles |Sewer, Artist, Crafter and Instructor|
+-----+-----+
```

(3)

Here is a brief explanation of the SQL query concerning the problem description:

- **All courses that are (a) related to 'interview':** As `like` is case-insensitive, we can use `%interview%` to match any string that contains 'interview', which can also include 'interviews'.
- **Sorted by course_rating in descending order and created in descending order (newest first):** order the data by `rating` in descending order and `created` in descending order.
- **Course rating should be firstly rounded to one decimal place:** use `round` to round the column `rating` to one decimal place.

```
1 course.createOrReplaceTempView("course")
2 spark.sql("select course_title as course,round(rating,1) as rating , created
from course where course_title like '%interview%' order by rating desc, created
desc").show(5,truncate=False)
```

```
+-----+-----+-----+
|course|rating|created|
+-----+-----+-----+
|Réaliser des interviews au rendu professionnel (PARTIE 2)|4.9|2022-08-12 14:54:06|
|Win your Product Management job interview with Big Tech's PM|4.8|2022-08-26 10:43:53|
|Get your Java dream job! Beginners interview preparation|4.8|2017-03-25 22:54:38|
|Angular interview questions with answers|4.6|2020-05-02 06:13:45|
|Software Testing Interview Masterclass: Ace the QA interview|4.6|2019-12-14 19:54:00|
+-----+-----+-----+
only showing top 5 rows
```